



HOMework 2 – SAMPLE SOLUTIONS

IMPORTANT NOTE

These homework solutions show multiple approaches and some optional extensions for most of the questions in the assignment. You don't need to submit all this in your assignments; they're included here just to help you learn more – because remember, the main goal of the homework assignments, and of the entire course, is to help you learn as much as you can, and develop your analytics skills as much as possible!

Question 3.1

Using the same data set (*credit_card_data.txt* or *credit_card_data-headers.txt*) as in Question 2.2, use the *ksvm* or *kknn* function to find a good classifier:

- (a) using cross-validation (do this for the *k*-nearest-neighbors model; SVM is optional); and
- (b) splitting the data into training, validation, and test data sets (pick either KNN or SVM; the other is optional).

SOLUTIONS:

(a)

There are different ways to do this. Three different methods are shown in `solution 3.1-a.R`. Just having one method is fine for your homework solutions. All three are shown below, for learning purposes. Another optional component shown below is using cross-validation for *ksvm*; this too did not need to be included in your solutions.

METHOD 1

The simplest approach, using *kknn*'s built-in cross-validation, is fine as a solution. `train.kknn` uses leave-one-out cross-validation, which sounds like a different type of cross-validation that I didn't mention in the videos – but if you watched the videos, you know it implicitly already! For each data point, it fits a model to *all* the other data points, and uses the remaining data point as a test – in other

words, if n is the number of data points, then leave-one-out cross-validation is the same as n -fold cross-validation.

Using this approach here are the results (using scaled data):

k	Correct	Percent correct	k	Correct	Percent correct
1,2,3,4	533	81.50%	18	557	85.17%
5	557	85.17%	19-20	556	85.02%
6	553	84.56%	21	555	84.86%
7	554	84.71%	22	554	84.71%
8	555	84.86%	23	552	84.40%
9	554	84.71%	24-25	553	84.56%
10-11	557	85.17%	26	552	84.40%
12	558	85.32%	27	550	84.10%
13-14	557	85.17%	28	548	83.79%
15-17	558	85.32%	29	549	83.94%
			30	550	84.10%

As before $k < 5$ is clearly worse than the rest, and value of k between 10 and 18 seem to do best. For unscaled data, the results are significantly worse (not shown here, but generally between 66% and 71%).

Note that technically, these runs just let us choose a model from among $k=1$ through $k=30$, but because there might be random effects in validation, to find an estimate of the model quality we'd have to run it on some test data that we didn't use for training/cross-validation.

METHOD 2

Some of you used the `cv.kknn` function in the `kknn` library. This approach is also shown in [solution 3.1-a.R](#).

METHOD 3

And others of you found the `caret` package in R that has the capability to run k -fold cross-validation (among other things). The built in functionality of the `caret` package gives ease of use but also the flexibility to tune different parameters and run different models. It's worth trying. This approach is also shown in [solution 3.1-a.R](#).

The main line of code is:

```
knn_fit <- train(as.factor(V11)~V1+V2+V3+V4+V5+V6+V7+V8+V9+V10,
```

```

data,
method = "knn", # choose knn model
trControl=trainControl(
  method="repeatedcv", # k-fold cross validation
  number=10, # number of folds (k in cross validation)
  repeats=5), # number of times to repeat k-fold cross
  validation
preProcess = c("center", "scale"), # standardize the data
tuneLength = kmax) # max number of neighbors (k in nearest
  neighbor)

```

The `trainControl` method allows us to determine the number of resampling iterations (“number”) and the number of folds to perform (“repeats”). The `train` function finally trains the model while allowing us to preprocess the data (scale and center) as well as select the number of k values to choose from.

ksvm Cross Validation

If you also tried cross-validation with `ksvm` (you didn’t need to), you could do that by including “`cross=k`” for k-fold cross-validation – for example, “`cross=10`” gives 10-fold cross-validation.

In the R code for Question 2.2 Part 1, you would replace the line

```

model <- ksvm(as.matrix(data[,1:10]),as.factor(data[,11]),type = "C-
svc",kernel = "vanilladot",C = 100,scaled = TRUE)

```

with

```

model <- ksvm(as.matrix(data[,1:10]),as.factor(data[,11]),type = "C-
svc",kernel = "vanilladot",C = 100,scaled = TRUE,cross=10)

```

`model@cross` shows the error measured by cross-validation, so `1-model@cross` is the estimate of the model’s fraction of points correctly classified; instead of the 86.4% correct classifications found in Question 2.2 Part 1 using scaled data, the cross-validated estimate is a little bit lower: 86.2%. That’s a difference of only about 1 correct prediction out of 654, so it’s not a big difference – meaning our initial model is a good one, and doesn’t seem to have been over-fit.

To compare models with different values of C, we can use that modification in the code in `solution 2.2-1.R`.

The results with scaled data to show that for `C=0.00001` or `C=0.00001`, only about 55% of points are classified correctly. At `C=0.001`, about 83% are classified correctly. At 0.01 and higher, the model achieves the 86.2% classification correctness we got above – a wide range of values of C gives a good

model. With unscaled data, just as before, finding a value of C to give a good model is harder.

<i>C</i>	<i>Percent correct (scaled data)</i>	<i>Percent correct (unscaled data)</i>
0.00001	54.76%	66.19%
0.0001	54.74%	68.50%
0.001	82.86%	75.38%
0.01	86.23%	81.50%
0.1	86.24%	85.48%
1	86.24%	78.89%
10	86.26%	58.81%
100	86.25%	70.04%
1000	86.25%	65.36%

(b)

As usual, there are lots of possible answers. File `solution 3.1-b.R` contains one approach.

In this approach, we first split the data into training, validation, and test sets. We used 60%, 20%, and 20%, but other splits are fine too as long as training has at least half.

Then, we fit 9 SVM models and 20 k-nearest-neighbor models to the training data, and evaluated them on the validation data.

We report the SVM model that does best in validation, and the KNN model that does best in validation. The code chose $C=0.01$ as the best SVM model, though it was equal with $C=0.1$, 1, 10, 100, and 1000, so any of them could've been chosen. The best KNN model was with $k=16$.

Then, we have an if statement that checks to see which model – the best SVM model or the best KNN model – performed best on the validation data. Whichever one it is, is the model we suggest using (and we report its performance on the test set).

Important note: In our code, the best SVM model ($C=0.01$) performs best on the validation data... but the best KNN model performs best on the test data. It might be tempting to therefore say, “Oh, let’s use the best KNN model.” Don’t give in to this temptation! If you do, you’re losing the value of separating validation and test sets. You’d essentially be using the test set to pick the best model, and then you’d (incorrectly) be using that same test set to estimate its quality – the selection bias from the validation step will be incorrectly included in the quality estimate.

You could’ve used a different approach – for example, only testing SVM models or only testing KNN models.

Some people also went beyond what we’ve covered, and tested models with different kernels – that’s also a good idea, and it’s possible to get better models that way.

Question 4.1

Describe a situation or problem from your job, everyday life, current events, etc., for which a clustering model would be appropriate. List some (up to 5) predictors that you might use.

Here's one answer.

An investor who wants to diversify a portfolio might want to cluster stocks, and then make sure the portfolio does not have too much money invested in any particular cluster.

A common way of clustering is to just classify each company by economic sector or size, but there might be deeper similarities that aren't captured by those factors. So, the investor might create factors related to each stock's performance (such as percent increase/decrease in price) in each quarter over the past 5 years, or each stock's performance in certain key days or intervals, etc. Stocks that behaved similarly would be clustered together.

Question 4.2

The iris data set `iris.txt` contains 150 data points, each with four predictor variables and one categorical response. The predictors are the width and length of the sepal and petal of flowers and the response is the type of flower. The data is available from the R library datasets and can be accessed with `iris` once the library is loaded. It is also available at the UCI Machine Learning Repository (<https://archive.ics.uci.edu/ml/datasets/Iris>). The response values are only given to see how well a specific method performed and should not be used to build the model.

Use the R function `kmeans` to cluster the points as well as possible. Report the best combination of predictors, your suggested value of k , and how well your best clustering predicts flower type.

Here's one possible solution. Please note that a good solution doesn't have to try all of the possibilities in the code; they're shown to help you learn, but they're not necessary.

The R code in file `solution_4.2.R` shows clustering solutions for $k=2,3,4,5$ using all factors, for both unscaled and scaled data.

	Unscaled data				Scaled data			
	Cluster	Setosa	Versi-color	Virgin-ica	Cluster	Setosa	Versi-color	Virgin-ica
k=2	1	50	3	0	1	50	0	0
	2	0	47	50	2	0	50	50
k=3	1	50	0	0	1	50	0	0
	2	0	48	14	2	0	47	14
	3	0	2	36	3	0	3	36
k=4	1	50	0	0	1	50	0	0
	2	0	27	1	2	0	27	2
	3	0	0	32	3	0	0	29
	4	0	23	17	4	0	23	19

k=5	1	50	0	0	1	28	0	0
	2	0	24	1	2	22	0	0
	3	0	0	24	3	0	27	2
	4	0	0	12	4	0	0	29
	5	0	26	13	5	0	23	19

Table 1. Results using all factors

For k=2, the setosa species is almost perfectly in one cluster, and the other two species (versicolor and virginica) are in the other cluster. For k=3,4,5, setosa is a perfect cluster. When k=4,5 there's a nice cluster of versicolor, a nice cluster or two of virginica, and a cluster of about 40 points that is mixed between the two. k=3 is a little more ambiguous – so even though there are 3 species, it turns out that k=4,5 work better.

The R code also shows clustering solutions for k=2,3,4,5 using only the Petal Length and Petal Width factors, for both unscaled and scaled data.

	Unscaled data				Scaled data			
	Cluster	Setosa	Versicolor	Virginica	Cluster	Setosa	Versicolor	Virginica
k=2	1	50	1	0	1	50	0	0
	2	0	49	50	2	0	50	50
k=3	1	50	0	0	1	50	0	0
	2	0	48	4	2	0	48	4
	3	0	2	46	3	0	2	46
k=4	1	50	0	0	1	50	0	0
	2	0	26	0	2	0	42	0
	3	0	0	35	3	0	0	27
	4	0	24	15	4	0	8	23
k=5	1	50	0	0	1	50	0	0
	2	0	22	0	2	0	23	0
	3	0	0	30	3	0	25	4
	4	0	0	13	4	0	0	25
	5	0	28	7	5	0	2	21

Table 2. Results using only Petal Length and Petal Width factors

Using only the Petal Length and Petal Width factors significantly improves the k=3 solution, and the k=5 solution. Notice that for k=4 especially, using scaled data is a big improvement over using unscaled data.

The R code also introduces the ggplot2 library for plotting, just for your learning pleasure – it's not required for the assignment.

Of course, we can only create the tables above because we happen to know the correct species for each data point. Normally when we're doing clustering, we don't have that information. Instead, we can look at a measure like the total distance between points and their cluster centers in each clustering solution, as shown in the elbow diagram below for scaled data using only the petal factors.




Figure 1. Elbow diagram for scaled data using only petal factors.

Based on this figure, the 3-cluster solution might be the one we would recommend, since $k=3$ is where the improvements level out.