

In [1]:

```
#reading the data
data <- read.table('./uscrime.txt', stringsAsFactors = F, header = T)
```

In [2]:

```
#printing the head of the data
head(data)
```

M	So	Ed	Po1	Po2	LF	M.F	Pop	NW	U1	U2	Wealth	Ineq	Prob	Cr
15.1	1	9.1	5.8	5.6	0.510	95.0	33	30.1	0.108	4.1	3940	26.1	0.084602	2.97
14.3	0	11.3	10.3	9.5	0.583	101.2	13	10.2	0.096	3.6	5570	19.4	0.029599	1.91
14.2	1	8.9	4.5	4.4	0.533	96.9	18	21.9	0.094	3.3	3180	25.0	0.083401	2.68
13.6	0	12.1	14.9	14.1	0.577	99.4	157	8.0	0.102	3.9	6730	16.7	0.015801	1.14
14.1	0	12.1	10.9	10.1	0.591	98.5	18	3.0	0.091	2.0	5780	17.4	0.041399	1.56
12.1	0	11.0	11.8	11.5	0.547	96.4	25	4.4	0.084	2.9	6890	12.6	0.034201	1.30

In [3]:

```
#scaling the data

#keeping the So variable which is a factor out of scaling. Also, keeping the dependent variable out.

drops <- c("So","Crime")
df <- data[ , !(names(data) %in% drops)]
scaled_data <- scale(df)

#binding the data altogether.
scaled_data <- cbind(scaled_data, data[,drops])
```

In [4]:

```
head(scaled_data)
```

M	Ed	Po1	Po2	LF	M.F	Pop
0.9886930	-1.3085099	-0.9085105	-0.8666988	-1.2667456	-1.12060499	-0.09500679
0.3521372	0.6580587	0.6056737	0.5280852	0.5396568	0.98341752	-0.62033844
0.2725678	-1.4872888	-1.3459415	-1.2958632	-0.6976051	-0.47582390	-0.48900552
-0.2048491	1.3731746	2.1535064	2.1732150	0.3911854	0.37257228	3.16204944
0.1929983	1.3731746	0.8075649	0.7426673	0.7376187	0.06714965	-0.48900552
-1.3983912	0.3898903	1.1104017	1.2433590	-0.3511718	-0.64550313	-0.30513945

In [5]:

```
options(warn=-1)
library(caret)
library(leaps)
library(MASS)
library(glmnet)
```

```
Loading required package: lattice
Loading required package: ggplot2
Loading required package: Matrix
Loading required package: foreach
Loaded glmnet 2.0-13
```

In [6]:

```
set.seed(1)
```

References:

<http://www.sthda.com/english/articles/37-model-selection-essentials-in-r/154-stepwise-regression-essentials-in-r/> (<http://www.sthda.com/english/articles/37-model-selection-essentials-in-r/154-stepwise-regression-essentials-in-r/>)

In [7]:

```
#defining a control method to do repeated cv

control <- trainControl(method = "repeatedcv", number = 5, repeats = 5)

#here we use the train function from caret package. method = "leapSeq" is for stepwise regression.
#tuneGrid = data.frame(nvmax = 1:15) means that it will try all the models with 1 features upto 15 features.
#the best 1-variable model, the best 2-variables model, ..., the best 15-variables model

lm_step <- train(Crime ~., data = scaled_data, method = "leapSeq",
                 tuneGrid = data.frame(nvmax = 1:15), trControl = control)
```

In [8]:

```
#here we print the stats for all nvmax best models

lm_step$results
```

nvmax	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
1	282.9300	0.4999792	226.2151	70.67819	0.2125575	53.74191
2	310.7000	0.4334649	248.3198	86.46990	0.2741488	72.16826
3	273.1027	0.5310749	214.0348	74.52462	0.2073685	61.93288
4	298.4097	0.4655277	236.8706	53.48436	0.1921420	42.46384
5	289.2306	0.4869785	233.2909	62.55661	0.2145011	56.34413
6	279.3289	0.5446330	221.5333	80.05025	0.2078361	68.68278
7	286.6147	0.5283167	224.7424	71.09796	0.1956827	64.51009
8	286.9273	0.4869795	227.3073	60.09801	0.2087288	54.08324
9	299.6539	0.4934172	231.2610	63.50853	0.1798948	59.63639
10	292.4513	0.5150836	228.5901	66.06529	0.1939786	55.64956
11	297.3778	0.4640345	231.5788	60.65984	0.1975678	57.45597
12	285.4662	0.5157689	224.6591	60.03197	0.2114072	57.48195
13	282.2549	0.5349635	222.1849	61.77533	0.2103352	53.76748
14	286.1901	0.5362130	224.7453	64.94866	0.1970353	57.80696
15	292.0083	0.5197076	229.6161	61.32444	0.2019892	56.76443

```
#it can be seen that the best model is one with 3 features. Also, good point to
note
# different seed values give different nvmax. When I set seed = 42, I got nvmax
= 6
lm_step$bestTune
```

	nvmax
3	3

```
#printing the summary to know the 3 features

summary(lm_step$finalModel)
```

	Forced in	Forced out
M	FALSE	FALSE
Ed	FALSE	FALSE
Po1	FALSE	FALSE
Po2	FALSE	FALSE
LF	FALSE	FALSE
M.F	FALSE	FALSE
Pop	FALSE	FALSE
NW	FALSE	FALSE
U1	FALSE	FALSE
U2	FALSE	FALSE
Wealth	FALSE	FALSE
Ineq	FALSE	FALSE
Prob	FALSE	FALSE
Time	FALSE	FALSE
So	FALSE	FALSE

Selection Algorithm: 'sequential replacement'

[illegible]

```
# * indicates that the features was included in the model.
# so the features selected are Ed, Pol, Ineq
#Using those features to fit the linear regression model
```

In [12]:

```
#fitting simple linear regression model on those 3 features
```

```
mod <- lm(Crime ~ Ed + Pol + Ineq, data = scaled_data)
summary(mod)
```

Call:

```
lm(formula = Crime ~ Ed + Pol + Ineq, data = scaled_data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-590.30	-102.06	-1.73	129.16	511.60

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	905.09	33.74	26.825	< 2e-16	***
Ed	176.61	53.32	3.312	0.00188	**
Pol	369.45	43.94	8.408	1.26e-10	***
Ineq	299.45	60.15	4.978	1.09e-05	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 231.3 on 43 degrees of freedom

Multiple R-squared: 0.6656, Adjusted R-squared: 0.6423

F-statistic: 28.53 on 3 and 43 DF, p-value: 2.59e-10

In [13]:

```
#all the coefficients are significant as per the p-value.
```

In [14]:

```
#now lets perform leave out one cross validation and get the R squared.
```

In [15]:

```
test <- numeric()
for (i in 1:nrow(scaled_data)){
  model <- lm(Crime ~ Ed + Pol +Ineq, data = scaled_data[-i,])
  test <- cbind(test, predict(model, newdata = scaled_data[i,]))
}
```

In [16]:

```
#function to calculate r-squared

compute_rsquared <- function(y_hat, y){
SSR <- sum((y_hat - y)^2)
SST <- sum((y - mean(y))^2)
rsquared <- 1 - SSR/SST
return (rsquared)
}
```

In [17]:

```
compute_rsquared(test, scaled_data$Crime)
```

0.574751286403875

In [18]:

```
#Using just three features we were able to get Rsquared of 0.57. Please note --
adding more features will
#definitely increase Rsquared. But we need to strike a balance between simplicit
y and generalization.
#i was able to get a pretty simple model using just three features.
```

In [19]:

```
full.model <- lm(Crime ~., data = scaled_data)
# Stepwise regression model
step.model <- stepAIC(full.model, direction = "both",
                      trace = FALSE)
summary(step.model)
```

```
Call:
lm(formula = Crime ~ M + Ed + Pol + M.F + U1 + U2 + Ineq + Prob,
    data = scaled_data)
```

Residuals:

Min	1Q	Median	3Q	Max
-444.70	-111.07	3.03	122.15	483.30

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	905.09	28.52	31.731	< 2e-16	***
M	117.28	42.10	2.786	0.00828	**
Ed	201.50	59.02	3.414	0.00153	**
Pol	305.07	46.14	6.613	8.26e-08	***
M.F	65.83	40.08	1.642	0.10874	
U1	-109.73	60.20	-1.823	0.07622	.
U2	158.22	61.22	2.585	0.01371	*
Ineq	244.70	55.69	4.394	8.63e-05	***
Prob	-86.31	33.89	-2.547	0.01505	*

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 195.5 on 38 degrees of freedom

Multiple R-squared: 0.7888, Adjusted R-squared: 0.7444

F-statistic: 17.74 on 8 and 38 DF, p-value: 1.159e-10

In [20]:

```
#So the above model selected 8 features and was able to get a rsquared of 0.7888
#and adjusted r-squared of 0.7444.
```

In [21]:

```
#using the above model... lets cross validate
test1 <- numeric()
for (i in 1:nrow(scaled_data)){
    model <- lm(Crime ~ M + Ed + Pol + M.F + U1 + U2 + Ineq + Prob, data =
scaled_data[-i,])
    test1 <- cbind(test1, predict(model, newdata = scaled_data[i,]))
}
```

In [22]:

```
compute_rsquared(test1, scaled_data$Crime)
```

0.667620969502124

In [23]:

```
#But it can be seen that M.F and U1 are not that significant  
#as per p-value so lets take that out.  
  
#using the above model... lets cross validate  
test2 <- numeric()  
for (i in 1:nrow(scaled_data)){  
  model <- lm(Crime ~ M + Ed + Pol + U2 + Ineq + Prob, data = scaled_data[-i,]  
)  
  test2 <- cbind(test2, predict(model, newdata = scaled_data[i,]))  
}
```

In [24]:

```
compute_rsquared(test2, scaled_data$Crime)
```

0.666163842867471

In [25]:

```
#So it can be seen that, we were able to get about the same r-squared by omittin  
g M.F and U1.  
#Occam's razor -- the simpler the better.  
#Hence we would go with the latter one with just 6 features.
```

Now it is quite a judgement call here. We have two models one with 3 features and other with 6 features. The final selection totally depends upon whether we want simplicity or generalization.

Lasso Regression

In [32]:

```
lasso_model = cv.glmnet(x = as.matrix(scaled_data[,-16]),  
                        y = as.matrix(scaled_data$Crime),  
                        alpha = 1, nfolds = 5,  
                        type.measure = "mse", family = "gaussian", standardize =  
F)
```

In [33]:

```
lasso_model
```

```
$lambda  
 [1] 260.2814612 237.1587736 216.0902418 196.8933803 179.4019150 163  
.4643433  
 [7] 148.9426216 135.7109696 123.6547811 112.6696312 102.6603717 93  
.5403072  
[13] 85.2304441 77.6588063 70.7598120 64.4737054 58.7460391 53
```



```
.5272029
[19] 48.7719937 44.4392242 40.4913660 36.8942246 33.6166434 30
.6302335
[25] 27.9091279 25.4297579 23.1706483 21.1122318 19.2366793 17
.5277457
[31] 15.9706291 14.5518424 13.2590969 12.0811952 11.0079352 10
.0300205
[37] 9.1389812 8.3270993 7.5873427 6.9133041 6.2991452 5
.7395465
[43] 5.2296610 4.7650723 4.3417565 3.9560468 3.6046025 3
.2843795
[49] 2.9926043 2.7267496 2.4845127 2.2637954 2.0626861 1
.8794427
[55] 1.7124782 1.5603464 1.4217295 1.2954270 1.1803448 1
.0754862
[61] 0.9799430 0.8928876 0.8135659 0.7412909 0.6754367 0
.6154328
[67] 0.5607594 0.5109431 0.4655523 0.4241939 0.3865097 0
.3521733
[73] 0.3208872 0.2923804 0.2664061
```

```
$cvm
 [1] 145826.80 142162.52 136720.36 128986.49 122615.74 117372.22 113
030.77
 [8] 109190.33 105894.32 103045.85 100968.76 99444.38 98316.42 97
599.12
[15] 96893.45 95717.79 93560.54 90268.30 87180.21 84856.90 82
851.73
[22] 81247.48 79741.87 78108.00 75845.12 73630.84 71379.19 68
963.85
[29] 67177.64 65867.71 64735.74 63927.27 63312.68 62864.51 62
936.45
[36] 63170.13 63362.86 63288.38 63349.75 63406.67 63468.32 63
593.58
[43] 63757.60 63946.58 64161.48 64412.40 64726.41 65256.90 65
993.48
[50] 66685.82 67208.62 67754.40 68262.80 68796.78 69325.67 69
780.42
[57] 70227.05 70660.63 71069.55 71459.47 71815.03 72174.90 72
513.98
[64] 72842.40 73135.52 73405.11 73658.93 73885.81 74099.19 74
296.06
[71] 74478.23 74629.45 74772.41 74904.31 75025.78
```

```
$cvstd
 [1] 30833.26 32189.43 31448.90 28670.62 26213.00 24057.37 22171.84
20419.36
 [9] 18899.04 17610.39 16619.38 15841.90 15263.54 14922.87 14824.25
14835.13
[17] 14712.28 14635.46 14689.81 14856.27 15011.28 15155.07 15400.04
15622.36
[25] 15237.15 14875.81 14449.95 14082.18 13791.42 13556.77 13392.32
13228.64
```

[33] 13080.19 12988.30 13105.46 13257.97 13303.73 13298.18 13271.08
13253.50
[41] 13247.52 13236.88 13226.06 13213.93 13199.09 13177.26 13133.24
13091.20
[49] 13113.98 13142.55 13054.33 13032.83 12990.46 12968.60 12939.69
12939.30
[57] 12945.52 12942.86 12949.56 12954.77 12961.84 12981.55 13016.57
13065.06
[65] 13109.91 13151.54 13193.91 13239.68 13279.49 13315.02 13351.15
13378.89
[73] 13403.56 13429.20 13451.80

\$cvup

[1] 176660.07 174351.95 168169.26 157657.11 148828.74 141429.58 135
202.61
[8] 129609.69 124793.36 120656.24 117588.13 115286.28 113579.96 112
522.00
[15] 111717.70 110552.92 108272.82 104903.75 101870.02 99713.16 97
863.01
[22] 96402.55 95141.91 93730.37 91082.27 88506.66 85829.14 83
046.04
[29] 80969.06 79424.48 78128.06 77155.91 76392.87 75852.80 76
041.91
[36] 76428.10 76666.59 76586.56 76620.83 76660.17 76715.85 76
830.46
[43] 76983.66 77160.50 77360.57 77589.66 77859.64 78348.10 79
107.46
[50] 79828.37 80262.95 80787.23 81253.26 81765.39 82265.35 82
719.72
[57] 83172.57 83603.48 84019.12 84414.24 84776.86 85156.45 85
530.54
[64] 85907.46 86245.43 86556.66 86852.84 87125.49 87378.68 87
611.07
[71] 87829.37 88008.34 88175.97 88333.50 88477.58

\$cvlo

[1] 114993.54 109973.09 105271.47 100315.88 96402.73 93314.85 90
858.93
[8] 88770.97 86995.29 85435.46 84349.38 83602.48 83052.89 82
676.25
[15] 82069.20 80882.65 78848.26 75632.84 72490.40 70000.63 67
840.44
[22] 66092.41 64341.83 62485.64 60607.98 58755.03 56929.25 54
881.67
[29] 53386.23 52310.94 51343.41 50698.63 50232.49 49876.21 49
830.99
[36] 49912.16 50059.13 49990.20 50078.68 50153.18 50220.80 50
356.70
[43] 50531.55 50732.65 50962.39 51235.14 51593.17 52165.69 52
879.50
[50] 53543.27 54154.29 54721.57 55272.34 55828.18 56385.98 56
841.13
[57] 57281.53 57717.77 58119.99 58504.70 58853.19 59193.35 59

497.41
[64] 59777.33 60025.60 60253.57 60465.01 60646.12 60819.70 60981.04
[71] 61127.08 61250.57 61368.85 61475.11 61573.99

\$nzero
s0 s1 s2 s3 s4 s5 s6 s7 s8 s9 s10 s11 s12 s13 s14 s15 s16
s17 s18 s19
0 1 1 1 1 1 1 1 1 1 1 1 1 1 3 4 4
5 5 5
s20 s21 s22 s23 s24 s25 s26 s27 s28 s29 s30 s31 s32 s33 s34 s35 s36
s37 s38 s39
5 5 5 6 7 8 8 8 8 8 8 9 9 9 9 10 10
10 11 11
s40 s41 s42 s43 s44 s45 s46 s47 s48 s49 s50 s51 s52 s53 s54 s55 s56
s57 s58 s59
11 11 11 11 11 11 11 12 12 12 13 13 13 14 14 14 14
14 15 15
s60 s61 s62 s63 s64 s65 s66 s67 s68 s69 s70 s71 s72 s73 s74
15 15 15 15 15 15 15 15 15 15 15 15 14 14 14

\$name
mse
"Mean-Squared Error"

\$glmnet.fit

Call: glmnet(x = as.matrix(scaled_data[, -16]), y = as.matrix(scale
d_data\$Crime), alpha = 1, family = "gaussian", standardize = F)

	Df	%Dev	Lambda
[1,]	0	0.00000	260.30000
[2,]	1	0.08027	237.20000
[3,]	1	0.14690	216.10000
[4,]	1	0.20220	196.90000
[5,]	1	0.24820	179.40000
[6,]	1	0.28630	163.50000
[7,]	1	0.31800	148.90000
[8,]	1	0.34430	135.70000
[9,]	1	0.36610	123.70000
[10,]	1	0.38420	112.70000
[11,]	1	0.39920	102.70000
[12,]	1	0.41170	93.54000
[13,]	1	0.42210	85.23000
[14,]	1	0.43070	77.66000
[15,]	3	0.44240	70.76000
[16,]	4	0.45870	64.47000
[17,]	4	0.48700	58.75000
[18,]	5	0.52490	53.53000
[19,]	5	0.55650	48.77000
[20,]	5	0.58260	44.44000
[21,]	5	0.60430	40.49000
[22,]	5	0.62240	36.89000

[23,]	5	0.63730	33.62000
[24,]	6	0.64980	30.63000
[25,]	7	0.66700	27.91000
[26,]	8	0.68230	25.43000
[27,]	8	0.69750	23.17000
[28,]	8	0.71000	21.11000
[29,]	8	0.72040	19.24000
[30,]	8	0.72900	17.53000
[31,]	8	0.73610	15.97000
[32,]	9	0.74290	14.55000
[33,]	9	0.75080	13.26000
[34,]	9	0.75730	12.08000
[35,]	9	0.76270	11.01000
[36,]	10	0.76790	10.03000
[37,]	10	0.77220	9.13900
[38,]	10	0.77580	8.32700
[39,]	11	0.77930	7.58700
[40,]	11	0.78230	6.91300
[41,]	11	0.78480	6.29900
[42,]	11	0.78680	5.74000
[43,]	11	0.78850	5.23000
[44,]	11	0.79000	4.76500
[45,]	11	0.79110	4.34200
[46,]	11	0.79210	3.95600
[47,]	11	0.79290	3.60500
[48,]	12	0.79360	3.28400
[49,]	12	0.79420	2.99300
[50,]	12	0.79470	2.72700
[51,]	13	0.79510	2.48500
[52,]	13	0.79550	2.26400
[53,]	13	0.79580	2.06300
[54,]	14	0.79610	1.87900
[55,]	14	0.79630	1.71200
[56,]	14	0.79650	1.56000
[57,]	14	0.79670	1.42200
[58,]	14	0.79690	1.29500
[59,]	15	0.79710	1.18000
[60,]	15	0.79800	1.07500
[61,]	15	0.79880	0.97990
[62,]	15	0.79950	0.89290
[63,]	15	0.80010	0.81360
[64,]	15	0.80060	0.74130
[65,]	15	0.80100	0.67540
[66,]	15	0.80130	0.61540
[67,]	15	0.80160	0.56080
[68,]	15	0.80180	0.51090
[69,]	15	0.80200	0.46560
[70,]	15	0.80220	0.42420
[71,]	15	0.80230	0.38650
[72,]	15	0.80240	0.35220
[73,]	14	0.80250	0.32090
[74,]	14	0.80260	0.29240
[75,]	14	0.80270	0.26640

```
[76,] 14 0.80270 0.24270
[77,] 14 0.80280 0.22120
[78,] 14 0.80280 0.20150
[79,] 14 0.80290 0.18360
[80,] 14 0.80290 0.16730
[81,] 14 0.80290 0.15240
[82,] 14 0.80290 0.13890
[83,] 14 0.80300 0.12660
[84,] 14 0.80300 0.11530
[85,] 14 0.80300 0.10510
[86,] 14 0.80300 0.09574
[87,] 14 0.80300 0.08724
[88,] 14 0.80300 0.07949
```

```
$lambda.min
```

```
[1] 12.0812
```

```
$lambda.1se
```

```
[1] 27.90913
```

```
attr(,"class")
```

```
[1] "cv.glmnet"
```

```
In [34]:
```


```
coef(lasso_model, s = lasso_model$lambda.min)
```

```
16 x 1 sparse Matrix of class "dgCMatrix"
```

```
      1
(Intercept) 905.08511
M            85.09006
Ed           115.31161
Po1          309.89409
Po2           .
LF            .
M.F           50.46511
Pop           .
NW            10.57842
U1           -21.05244
U2            52.84583
Wealth       .
Ineq         182.41508
Prob        -75.82115
Time         .
So           .
```

In [35]:

```
plot(lasso_model)
```



In [36]:

```
#fitting a model with the above 9 variables. These variables are same to stepAIC  
  
lm_lasso <- lm(Crime ~ M + Ed + Po1 + M.F + NW + U1 + U2 + Ineq + Prob, data =  
scaled_data)
```

In [37]:

```
summary(lm_lasso)
```

Call:

```
lm(formula = Crime ~ M + Ed + Pol + M.F + NW + U1 + U2 + Ineq +  
    Prob, data = scaled_data)
```

Residuals:

Min	1Q	Median	3Q	Max
-439.2	-102.2	-6.3	124.1	476.6

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	905.09	28.87	31.352	< 2e-16	***
M	111.23	46.83	2.375	0.022820	*
Ed	203.63	60.12	3.387	0.001687	**
Pol	297.89	52.08	5.719	1.51e-06	***
M.F	68.74	41.63	1.651	0.107134	
NW	16.55	53.15	0.311	0.757222	
U1	-109.46	60.94	-1.796	0.080609	.
U2	156.94	62.09	2.528	0.015889	*
Ineq	236.70	61.95	3.821	0.000492	***
Prob	-89.99	36.28	-2.481	0.017791	*

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 197.9 on 37 degrees of freedom

Multiple R-squared: 0.7894, Adjusted R-squared: 0.7381

F-statistic: 15.41 on 9 and 37 DF, p-value: 4.881e-10

Point to note -- different value of seeds gives different results. One more point to note, the significant variables from the above model are same as the variables selected from step AIC. So I will not rerun the lm model on those 6 variables.

Elastic Net

In [102]:

```
#we loop over different values of alpha. In each loop we select, minimum cross v  
alidation error and lambda.min
```

```
mse_list <- numeric()  
find_alpha <- function(num, scaled_data){  
  alpha <- num  
  elastic_net <- cv.glmnet(x=as.matrix(scaled_data[,-16]),  
                           y=as.matrix(scaled_data[,16]),  
                           alpha = alpha,  
                           nfolds=5,  
                           type.measure="mse",  
                           family="gaussian",  
                           standardize=FALSE)  
  
  mse_list <-<- cbind(mse_list, c(alpha, min(elastic_net$cvm),elastic_net$lambda.min))  
}
```

In [103]:

```
#looping over different values of alpha
```

```
for (i in seq(.01,1,by = .01)){find_alpha(i,scaled_data)}
```

In [106]:

```
minIndex <- which.min(mse_list[2,])
```

In [107]:

```
mse_list[2,minIndex]
```

```
50367.424492338
```

In [108]:

```
mse_list[1, minIndex]
```

```
0.89
```


In [109]:

```
#Using alpha = 0.89 and running the elastic net
```

```
elastic_net_final <- cv.glmnet(x=as.matrix(scaled_data[,-16]),  
                              y=as.matrix(scaled_data[,16]),  
                              alpha = 0.89,  
                              nfolds=5,  
                              type.measure="mse",  
                              family="gaussian",  
                              standardize=FALSE)
```

In [110]:

```
coef(elastic_net_final, s = elastic_net_final$lambda.min)
```

16 x 1 sparse Matrix of class "dgCMatrix"

```
      1  
(Intercept) 901.772626  
M            110.002785  
Ed           186.022954  
Po1          292.041940  
Po2           .  
LF           -2.821441  
M.F          51.181343  
Pop         -25.993778  
NW           21.045777  
U1          -85.025620  
U2          131.958964  
Wealth       72.742142  
Ineq        272.634500  
Prob        -91.904515  
Time        -3.396470  
So           9.730411
```

In [111]:

```
#So elastic_net selects 14 variables. Let's rerun lm using those features.
```

In [113]:

```
lm_elastic <- lm(Crime ~ ., data = scaled_data[, -4])
```

In [114]:

```
summary(lm_elastic)
```

Call:

```
lm(formula = Crime ~ ., data = scaled_data[, -4])
```

Residuals:

Min	1Q	Median	3Q	Max
-442.55	-116.46	8.86	118.26	473.49

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	903.155	58.889	15.336	2.66e-16	***
M	112.934	52.244	2.162	0.038232	*
Ed	198.350	68.044	2.915	0.006445	**
Pol	286.864	71.091	4.035	0.000317	***
LF	-11.321	56.896	-0.199	0.843538	
M.F	53.684	59.798	0.898	0.376026	
Pop	-29.833	48.950	-0.609	0.546523	
NW	25.149	63.619	0.395	0.695239	
U1	-97.649	75.332	-1.296	0.204164	
U2	143.034	69.378	2.062	0.047441	*
Wealth	87.540	99.662	0.878	0.386292	
Ineq	290.076	90.023	3.222	0.002921	**
Prob	-97.432	49.655	-1.962	0.058484	.
Time	-7.991	47.425	-0.168	0.867251	
So	5.669	148.100	0.038	0.969705	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 208.6 on 32 degrees of freedom

Multiple R-squared: 0.7976, Adjusted R-squared: 0.709

F-statistic: 9.006 on 14 and 32 DF, p-value: 1.673e-07

In [115]:

```
#Again it can be seen that the significant variables are M, Ed, Pol, U2, Ineq, Prob. However lets perform #cv using those 14 variables.
```

In [116]:

```
test4 <- numeric()
for (i in 1:nrow(scaled_data)){
  model <- lm(Crime ~ ., data = scaled_data[-i,-4])
  test4 <- cbind(test4, predict(model, newdata = scaled_data[i,-4]))
}
```

In [117]:

```
compute_rsquared(test4, scaled_data$Crime)
```

0.493682047286474

In [118]:

```
#it can be seen that using those 14 variables we got a huge drop in r_square = 0.49
```

We saw that all models agreed on the final 6 variables - M, Ed, Po1, U2, Ineq, Prob. Also, I was able to get a very simple model with just 3 features which had decent performance considering that we have 3 features. I'm amazed that all other methods were selecting more features (more than 3) with very little improvement in performance. I would prefer as simple models as possible because it makes it easy to explain it to someone. Having worked as a data scientist, you always need to explain what the model is doing to the senior management.